Founder Decision Matrix

Polymarket Strategy Lane Selection — Data-Driven Assessment

Date: February 12, 2026 **Input:** Grok empirical research (Blocks 0–6) **Analyst:** Claude Opus **Decision maker:** Shay

---

Executive Summary

Grok returned hard data for ~35% of queried fields. The remaining 65% came back NO DATA. This is itself a major finding — most Polymarket strategy categories have zero publicly documented, verifiable performance data. The decision must be made under radical uncertainty for 3 of 5 categories.

Three findings that reshape the landscape:

1. **Fee structure is materially different from our prior model.** Polymarket docs confirm zero winner/redemption fee on standard markets. If true, NegRisk arb cost drops from ~3.5% to gas-only (~0.1–0.5%). This is the single most important data point in the entire return.
2. **Mechanical arb is the only category with documented, quantified profit data** ($29M NegRisk + $10.6M binary = $39.6M total, April 2024–April 2025). Every other category has zero verifiable P&L data.
3. **Market making rebates are real, expanding, and currently 100% redistribution** on fee-enabled markets. This is a documented, structural revenue stream — not speculative edge.

The data gap problem: Domain Specialization, Behavioral/Flow, and Model-Based/AI returned near-total NO DATA. This does not mean these strategies don't work — it means we cannot score them with evidence. Scoring them would be opinion, not analysis. They are classified CONDITIONAL pending dedicated research.

Section 1: Critical Data Findings

Before scoring, these findings from Grok's return require explicit acknowledgment because they change prior assumptions.

Finding 1: Fee Structure Revision

Fee Component	Our Prior Assumption	Grok Finding	Impact
Winner/redemption fee (standard markets)	2% on winning position	**Zero** (docs.polymarket.com, Feb 2026)	NegRisk cost model drops ~2%. Thresholds must be recalculated.

Fee Component	Our Prior Assumption	Grok Finding	Impact
Trading fee (standard markets)	Zero	Confirmed zero	No change.
15-min crypto taker fee	Existed	Confirmed (curve-based, max ~3% at 50¢)	Maker rebates funded by this.
NegRiskAdapter feeRate	Unknown (highest priority)	**Still unknown** — "configurable (0 or positive)"	STILL the single biggest unknown. Must be read on-chain.
Polymarket US taker fee	0.01% (prior reference)	**0.10%** (min $0.0010)	10x higher than prior reference. Corrected.
Gas (Polygon)	~$0.025/tx	$0.01–0.20/tx	Range wider than assumed. Convert-specific unknown.

Action required: If zero winner fee is confirmed, NegRisk arb viable threshold drops from $\Sigma \leq 0.965$ to approximately $\Sigma \leq 0.995$ (gas only). This is a transformative change — far more opportunities would qualify. Must verify by placing a $1 micro-trade and checking redemption.

Finding 2: Platform Scale

Metric	Value	Source
Monthly volume (Jan 2026)	$7.66B	Yahoo Finance
Active markets	25,085	DeFi Rate
Weekly active wallets	300K+ peaks	Dune
Profitable wallets	30% (revised up from 7.6%)	DeFi Oasis, Dec 2025
Polymarket US status	Live (DCM), $450M+ first month	Regulatory filings

The platform is 10x larger than when NegRisk arb data was collected (2024). More volume = more repricing events = potentially more arb opportunities, but also more bots.

Finding 3: Data Desert for Non-Arb Strategies

Category	Fields with data	Fields NO DATA	Data quality
Mechanical Arb	8 of 17	9 of 17	Moderate (IMDEA anchor)
Domain Specialization	0 of 12	12 of 12	**Zero**
Behavioral/Flow	0 of 10	10 of 10	**Zero**

Category	Fields with data	Fields NO DATA	Data quality
Model-Based/AI	1 of 12	11 of 12	Near-zero (Vitalik only)
Market Making	3 of 10	7 of 10	Low-moderate

Section 2: Candidate Scoring

Methodology: Scores derived from Grok data where available. Where NO DATA, scored conservatively (low) and flagged with ⚠. A score of "?" means no basis for any estimate — treated as 3/10 in totals (below midpoint, reflecting uncertainty penalty).

Scale: 1 = worst, 10 = best.

Full 15-Candidate Matrix

Dimension	1A NegRisk YES	1B NegRisk Convert	1C Binary Arb	1D Cross-Plat	2A Weather	2B Political	2C Niche	3A Whale	3B Overreaction	3C Post-Resol
Edge Durability	5	6	2	3⚠	?⚠	?⚠	?⚠	?⚠	?⚠	?⚠
Infra Barrier (10=easy)	7	5	3	4	6⚠	8⚠	8⚠	5⚠	6⚠	7⚠
Scalability	4	4	6	3⚠	2⚠	7⚠	1⚠	5⚠	3⚠	2⚠
Competition (10=low)	4	7	1	5⚠	7⚠	4⚠	8⚠	4⚠	6⚠	5⚠
Strategic Fit	7	8	2	4	5	6	4	5	4	3
Learning Compound	5	6	2	3	7⚠	8⚠	5⚠	6⚠	5⚠	3⚠
TOTAL (/60)	**32**	**36**	**16**	**22**	**30**	**36**	**29**	**28**	**27**	**23**
⚠ **count**	0	0	0	2	6	6	6	6	6	6

Score Justifications (Top Candidates Only)

1B NegRisk Convert — 36 (0 ⚠)

- Edge Durability (6): Convert flow is less documented than Buy-All-YES, suggesting lower bot competition. But IMDEA paper shows market is being studied — bots will follow.
- Competition (7): Highest score in mechanical arb. Grok returned zero documented convert-flow arb wallets — either unused or undiscovered. Both are good for us short-term.
- Strategic Fit (8): Highest in matrix. Systems-oriented (API + on-chain), rule-expressible, immediate capital return (no lockup), measurable.
- Learning Compound (6): Skills transfer to broader DeFi protocol interaction, on-chain mechanics, contract reading.

2B Political — 36 (6 ⚠️)

- Scores well on Strategic Fit (6) and Learning Compound (8) — political analysis skills compound across election cycles.
- But ALL scores are ⚠️ flagged — zero empirical data supports any of them. This is projection, not measurement.
- If scored at 3 for all ⚠️ fields: drops to 26/60.

4A LLM Probability — 35 (3 ⚠️)

- Vitalik's documented 16% return ($70K on $440K) is the only non-arb performance data in the entire dataset.
- But methodology was "fading long-shots" — this is a behavioral insight dressed as model-based trading.
- Learning Compound (8): Highest in category. LLM calibration skills compound across all market types and future AI applications.
- Competition (4): LLMs are commoditizing — every retail trader has access to Claude/GPT. Edge may be in prompt engineering and information retrieval, not model quality.

1A NegRisk YES — 32 (0 ⚠️)

- Only candidate with $29M documented profit data.
- Edge Durability (5): IMDEA data is from election supercycle. Grok confirms spread compression in 2026.
- Strategic Fit (7): Rule-based, testable, measurable — core Ziva.ai philosophy.
- Scalability (4): Limited by event universe size and resolution lockup.

5A Passive MM — 30 (0 ⚠️)

- Edge Durability (6): Maker rebates are structural (funded by taker fees), not speculative. Currently 100% redistribution.
- Infrastructure Barrier (4): Requires real-time quote management, post-only order support, position tracking.

- Competition (4): ICE institutional entry is documented and will compress spreads.

Section 3: GO / CONDITIONAL / KILL

KILL — Do Not Pursue

Candidate	Verdict	Reason
1C Binary YES+NO Arb	KILL	Sub-second half-life confirmed by multiple sources. Pure HFT. $10.6M extracted by bots we cannot compete with. Zero evidence of accessibility to non-HFT.
5B Active MM	KILL	Capital requirement ($50K+), infrastructure complexity (hedging, inventory management), and ICE institutional competition make this inaccessible at our stage. No documented P&L from any non-institutional maker.
3C Post-Resolution	KILL	Niche, infrequent, thin margin, and likely already captured by bots monitoring resolution oracles. Zero documented returns.

GO — Data Supports Immediate V1 Testing

Candidate	Verdict	Reason	First Action
1A NegRisk YES	GO	Only strategy with $29M documented profit. Fee revision (zero winner fee) dramatically improves economics. Two-lane observation protocol already designed.	Verify zero winner fee via $1 micro-trade. Then deploy 14-day scanner.
1B NegRisk Convert	GO	Zero documented competition on convert flow. Potential structural cost advantage (no lockup, immediate capital return). Same scanner infrastructure as 1A.	Read feeRate on-chain from deployed NegRiskAdapter contract. If ≤ 0.5%, proceed alongside 1A.

Critical note: 1A and 1B share 90% of infrastructure (same scanner, same events, same spread computation). Testing both adds near-zero marginal cost. The two-lane observation protocol we already designed is the correct V1.

CONDITIONAL — Cannot Score Without Dedicated Research

Candidate	Verdict	Condition to Convert to GO	Condition to Convert to KILL
2B Political	CONDITIONAL	Dedicated research phase: (1) Calibration analysis of Polymarket political markets vs outcomes, (2) Identify ≥ 3 systematic biases, (3) Paper-trade 20 positions	Markets well-calibrated (no bias), or liquidity too thin for meaningful deployment, or political markets migrate to Polymarket US with restricted access
4A LLM Probability	CONDITIONAL	(1) Build calibration benchmark: LLM probability estimates vs market prices vs actual outcomes across 50+ resolved markets, (2) Demonstrate ≥ 3% edge after costs	LLM estimates no better than market consensus, or edge exists but only on markets too thin to trade
2A Weather	CONDITIONAL	(1) Count active weather markets (need ≥ 5), (2) Measure model-vs-market divergence on 10+ markets, (3) Confirm liquidity ≥ $500/market	Fewer than 5 weather markets, or model-market divergence < 3%, or liquidity < $200/market
4C Cross-Dep	CONDITIONAL	(1) Reproduce IMDEA dependency detection, (2) Identify ≥ 10 active dependent pairs with > 3% divergence	Detection accuracy < 70%, or divergence < 2%, or resolution timing misaligned by > 90 days
2C Niche Quant	CONDITIONAL	(1) Count active niche-quantifiable markets (need ≥ 3), (2) Verify base-rate predictability on 5+ markets	Fewer than 3 markets, or markets already well-priced
3A Whale Tracking	CONDITIONAL	(1) Build whale identification pipeline, (2) Measure convergence signal quality on 20+ historical events	> 50% obfuscation rate, or signal-to-noise ratio < 1.5:1
3B Overreaction	CONDITIONAL	(1) Historical analysis of > 20% intraday moves (frequency, reversion rate)	Sharp moves < 1/month, or reversion rate < 50%
1D Cross-Platform	CONDITIONAL	(1) Map market overlap between Polymarket Global and Kalshi, (2) Measure price divergence on ≥ 10 equivalent markets	Overlap < 20 markets, or settlement mismatch > 5%, or fees exceed typical spread
5A Passive MM	CONDITIONAL	(1) Document current rebate economics on fee-enabled markets, (2) Estimate adverse selection cost, (3) Calculate minimum capital for positive expected value	Rebates discontinued, or adverse selection > spread revenue, or minimum capital > $10K

Section 4: Strategic Classification

Classification	Candidates	Reasoning
Long-term asymmetric advantage	4A LLM Probability, 2B Political, 4C Cross-Market Dependency	Learning compounds. Infrastructure creates moat. Skills transfer across markets and even beyond Polymarket. But: zero current evidence of profitability. These are bets on capability building, not proven edges.
Short-term extraction	1A NegRisk YES, 1B NegRisk Convert, 3A Whale Tracking	Edge is real (1A documented) or plausible (1B, 3A) but structurally compressing. Extract value while building long-term capabilities. Use profits to fund capability development.
Structurally declining	1C Binary Arb, 5B Active MM	Fully captured by HFT bots and institutional market makers. No viable entry point for our constraints.

Section 5: Recommended Path

V1 Testing Lane: 1A + 1B Combined (NegRisk Two-Lane Observation)

Why this lane, why now, why not alternatives:

This is the only strategy category with documented, quantified, multi-million-dollar profit data. The fee structure revision (zero winner fee if confirmed) transforms the economics from marginal to potentially strong. Both lanes share 90% infrastructure — testing both costs nearly nothing incremental. And crucially, this is an *observation protocol*, not a trading system — zero capital risk while generating decision-grade data for every other lane decision.

V2 Preparation Lane: 4A LLM Probability Estimation

Why: Highest learning-compounding score in the dataset. Vitalik's 16% documented return (the only non-arb performance data anywhere) suggests the edge exists. LLM calibration infrastructure transfers to every future strategy — it's a platform capability, not a single play. Can be developed in parallel with V1 observation at low cost (API calls + prompt engineering).

Top Kill Candidates (Skip Entirely)

1C Binary YES+NO Arb, 5B Active Inventory-Managed MM, 3C Post-Resolution Arb.

What Would Change This Recommendation

Single most important data point: The NegRiskAdapter on-chain feeRate. If it's > 2%, Lane B loses its structural advantage and reverts to "same as Lane A." If it's 0%, Lane B becomes the dominant lane (faster

capital turnover, no lockup, potentially lower cost). This is a 10-minute on-chain query that hasn't been done yet.

Second most important: Confirmation of zero winner fee via actual trade. If Grok is wrong and there IS a winner fee, the entire NegRisk cost model reverts to our prior (tighter thresholds, fewer opportunities). This is a $1 experiment.

Section 6: Uncertainty Register

This matrix was built on 35% data and 65% void. The following uncertainties are ranked by decision impact.

Rank	Uncertainty	Impact if Wrong	Resolution Method	Cost	Time
1	NegRiskAdapter feeRate	Lane B viability pivots entirely	Read on-chain from Polygonscan	$0	10 min
2	Zero winner fee on standard markets	NegRisk threshold changes from 0.995 to 0.965	Place $1 trade, observe redemption	$1	1 day
3	Current NegRisk spread distribution	Determines if opportunities exist at all	14-day scanner (already designed)	$0 (compute)	14 days
4	Opportunity half-life in 2026	Determines if we can execute fast enough	Same scanner, 5s polling during events	$0	14 days
5	Domain spec / behavioral / AI profitability	Determines V2+ strategy direction	Dedicated research phases (separate prompts)	Time only	2–4 weeks each

---

Decision Checklist for Shay

Before authorizing V1 Two-Lane Observation:

- [ ] Review this matrix and challenge any scores
- [ ] Confirm comfort with GO classification for 1A + 1B
- [ ] Confirm comfort with KILL classification for 1C, 5B, 3C
- [ ] Authorize Fee Truth Check: (a) read feeRate on-chain, (b) place $1 micro-trade
- [ ] Authorize 14-day observation scanner build (zero capital risk)
- [ ] Decide: pursue V2 preparation (4A LLM) in parallel, or sequential after V1?

If ANY box is unchecked or challenged, we discuss before building.

End of Founder Decision Matrix. 15 candidates evaluated. 3 killed. 2 approved for testing. 10 conditional pending dedicated research. All scores data-sourced or uncertainty-flagged.